



**APRIL 3-4, 2025**  
MARINA BAY SANDS / SINGAPORE

#BHAS @BlackHatEvents



# KubeAPI-Inspector: discover the secrets hidden in apis

Qiqi Xu

#BHAS @BlackHatEvents



# About us

- Security researcher at the Dawn Security Lab of JD.com
- Dawn Security Lab established in March 2021, Dawn Security Lab focuses
  - on cutting-edge offensive and defensive security research and product development.
- <https://dawnslab.jd.com/>

# Agenda

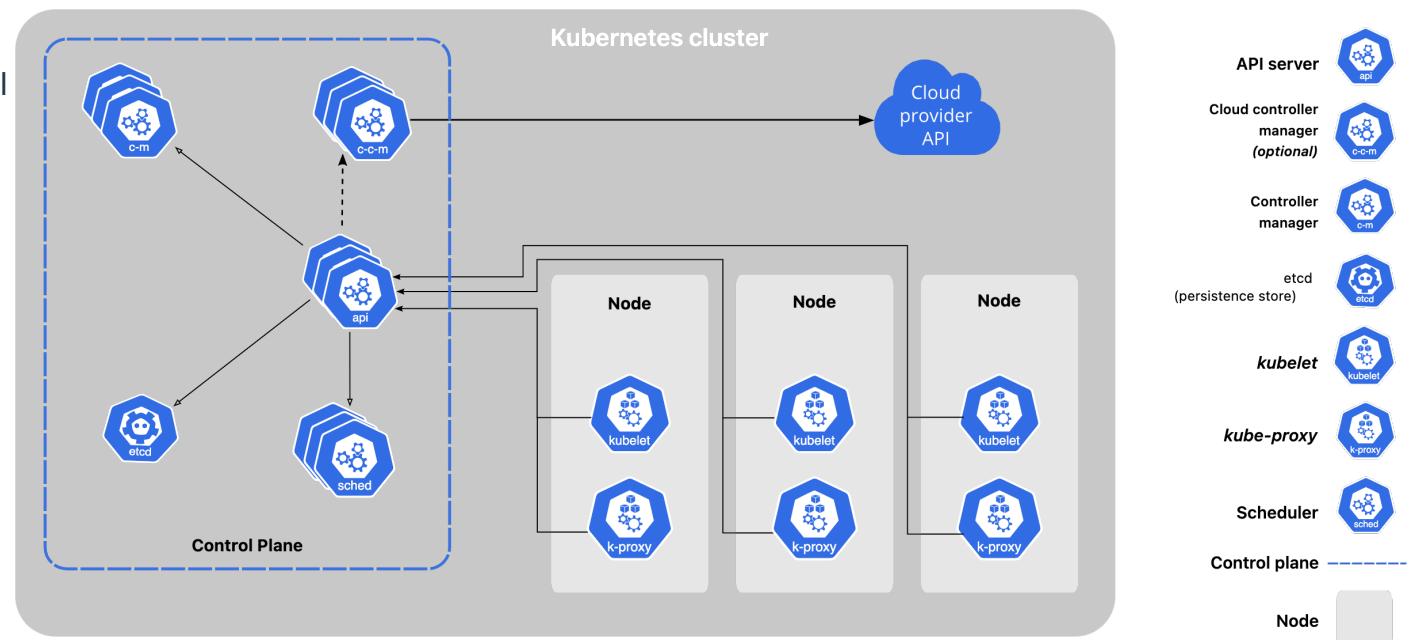
- Introduction
- How the tool works
- Workshop in action & Demo
- Next & Future



# Introduction

# In Kubernetes, API is "everything"

- External components and users all communicate with API server
- Everything operates(CRUD) via API
- Multiple scenarios: hybrid cloud, multi-tenancy, AI/ML

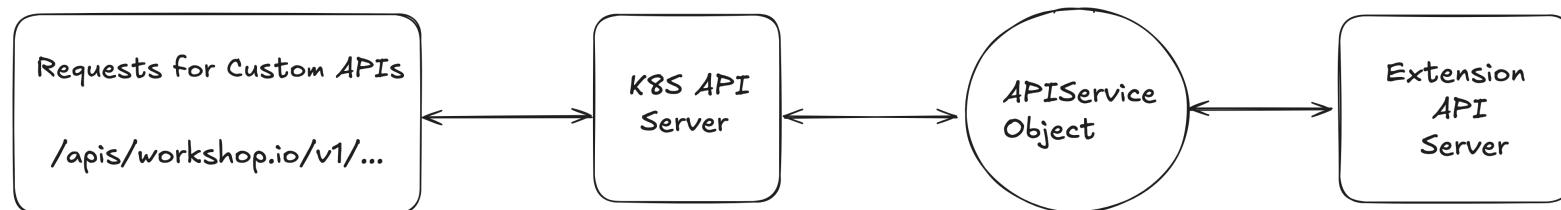


Ref: <https://kubernetes.io/docs/concepts/overview/>

#BHAS @BlackHatEvents

# The Extension API Server

- Extension API Server are a core extension capability of Kubernetes.
- Through Extension API Server, you can add functionality to Kubernetes.
- Independent API service



# Bridging the Gap in K8S API Security

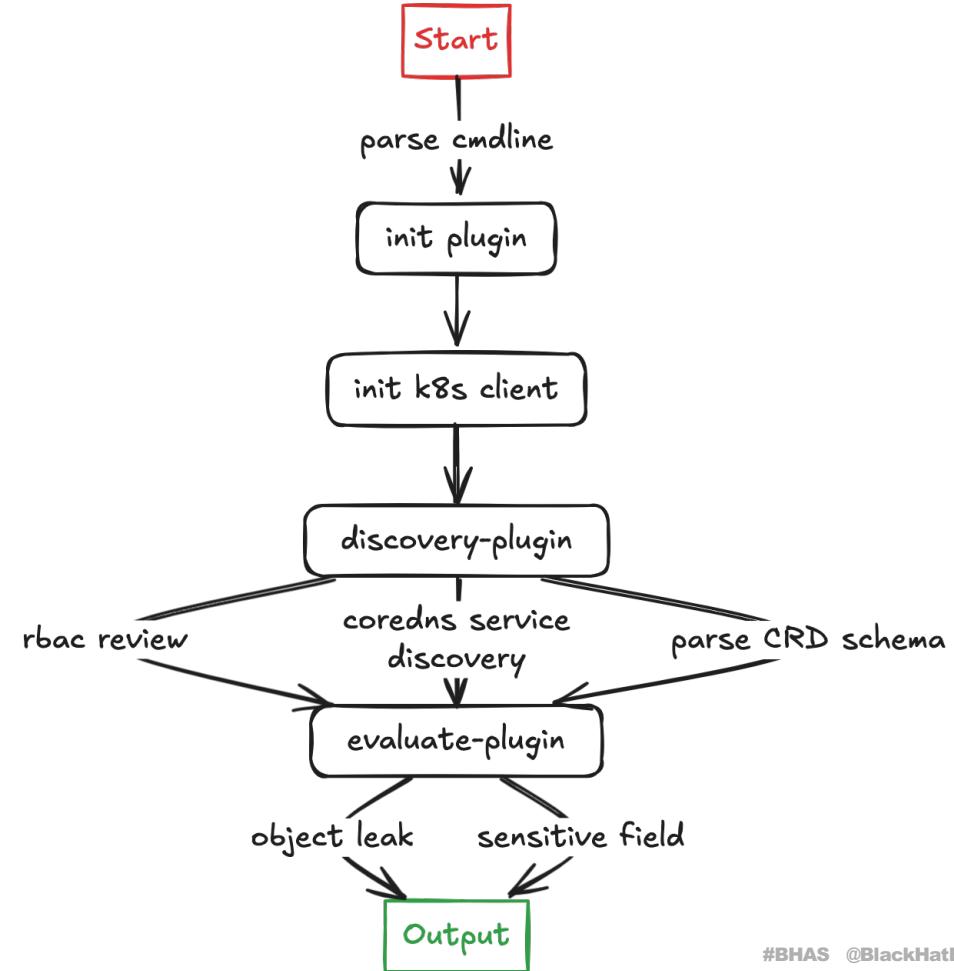
- Excellent open-source K8S security projects exist: CDK (penetration testing), Kubescape (misconfiguration scanning), Falco (runtime security), Cilium (network security)
- While these tools excel at container runtime and cluster baseline security, they offer limited Kubernetes API security analysis.
- We introduce kubeapi-inspector to fill this gap, focusing on identifying API-level vulnerabilities.



## How the tool works

# Feature

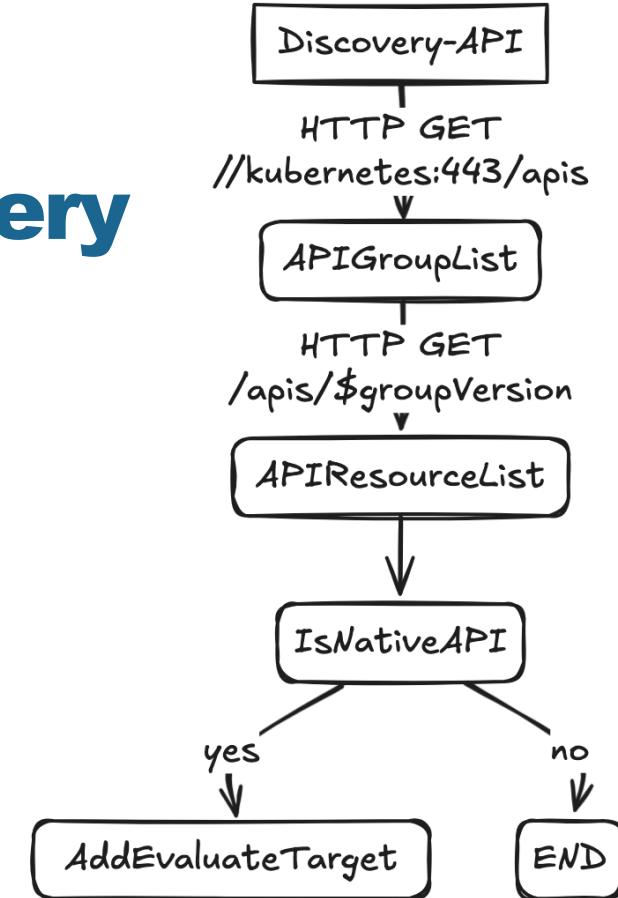
- Single Binary, Built with Golang
- Post-Exploitation tools
- Auto discover api token/address/endpoint
- Find object sensitive field
- api fuzzing



# Automated Target Discovery

- Use in-cluster config (K8S auto-mounts default service account).
- Use Client-go Discovery API to find K8S endpoints.

```
+ ~ kubectl --as=bh --as-group=system:authenticated auth can-i --list
Resources          Non-Resource URLs      Resource Names   Verbs
selfsubjectaccessreviews.authorization.k8s.io  []              []             [create]
selfsubjectrulesreviews.authorization.k8s.io    []              []             [create]
[]/api/*          []              []             [get]
[]/api            []              []             [get]
[]/apis/*        []              []             [get]
[]/apis           []              []             [get]
[]/healthz        []              []             [get]
[]/healthz        []              []             [get]
[]/livez          []              []             [get]
[]/livez          []              []             [get]
[]/openapi/*      []              []             [get]
[]/openapi         []              []             [get]
[]/readyz         []              []             [get]
[]/readyz         []              []             [get]
[]/version/*      []              []             [get]
[]/version         []              []             [get]
[]/version         []              []             [get]
[]/version         []              []             [get]
```



# Service discovery? Target discovery!

- Kubernetes uses CoreDNS as its default cluster DNS.
- CoreDNS versions before 1.9.0 have a Kubernetes plugin that can leak cluster info.
- This feature can help use to find other targets.

```
root@victim-db69bd6d7-7z8qm:/# dig srv *.*.*.svc.cluster.local  
;; <>> DiG 9.18.30-0ubuntu0.24.04.2-Ubuntu <>> srv *.*.*.svc.cluster.local  
;; global options: +cmd  
;; Got answer:  
;; WARNING: .local is reserved for Multicast DNS  
;; You are currently testing what happens when an mDNS query is leaked to DNS  
;; ->>HEADER<- opcode: QUERY, status: NOERROR, id: 56231  
;; flags: qr aa rd; QUERY: 1, ANSWER: 5, AUTHORITY: 0, ADDITIONAL: 5  
;; WARNING: recursion requested but not available  
  
;; OPT PSEUDORESECTION:  
;; EDNS: version: 0, flags:; udp: 1232  
;; COOKIE: 625f3e125a97288d (echoed)  
;; QUESTION SECTION:  
.*.*.*.svc.cluster.local. IN SRV  
  
;; ANSWER SECTION:  
.*.*.*.svc.cluster.local. 22 IN SRV 0 16 9153 172-17-0-2.kube-dns.kube-system.svc.cluster.local.  
.*.*.*.svc.cluster.local. 22 IN SRV 0 16 53 172-17-0-2.kube-dns.kube-system.svc.cluster.local.  
.*.*.*.svc.cluster.local. 22 IN SRV 0 16 8443 192-168-49-2.kubernetes.default.svc.cluster.local.  
.*.*.*.svc.cluster.local. 22 IN SRV 0 16 80 172-17-0-4.nginx.default.svc.cluster.local.  
.*.*.*.svc.cluster.local. 22 IN SRV 0 16 80 172-17-0-3.nginx.default.svc.cluster.local.
```

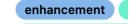
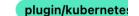
## Information Disclosure Issue: Add feature to disable wildcard lookups #4984

 Closed

micahhausler opened on Nov 16, 2021 · edited by micahhausler

Edits ⋮

Assignees  
No one assigned

Labels  
 enhancement  
 plugin/kubernetes

Opening this as a public issue, as there are already [public discussions](#) and [integration into penetration testing tooling](#).

CoreDNS's wildcard feature (specifically [wildcards in Kubernetes](#)) present an information disclosure.

CoreDNS has several plugins that support wildcard lookups, including the `file`, `etcd`, and `kubernetes` plugins. In environments such as Kubernetes where Service Discovery is desired, full enumeration of Services, Namespaces, and Endpoint IPs is possible even when Kubernetes API permissions or authentication is denied.

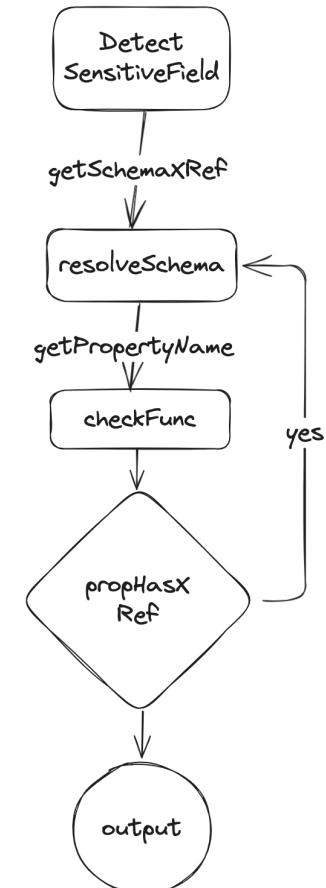
Ref: <https://github.com/coredns/coredns/issues/4984>

#BHAS @BlackHatEvents

# Detecting Sensitive API Fields

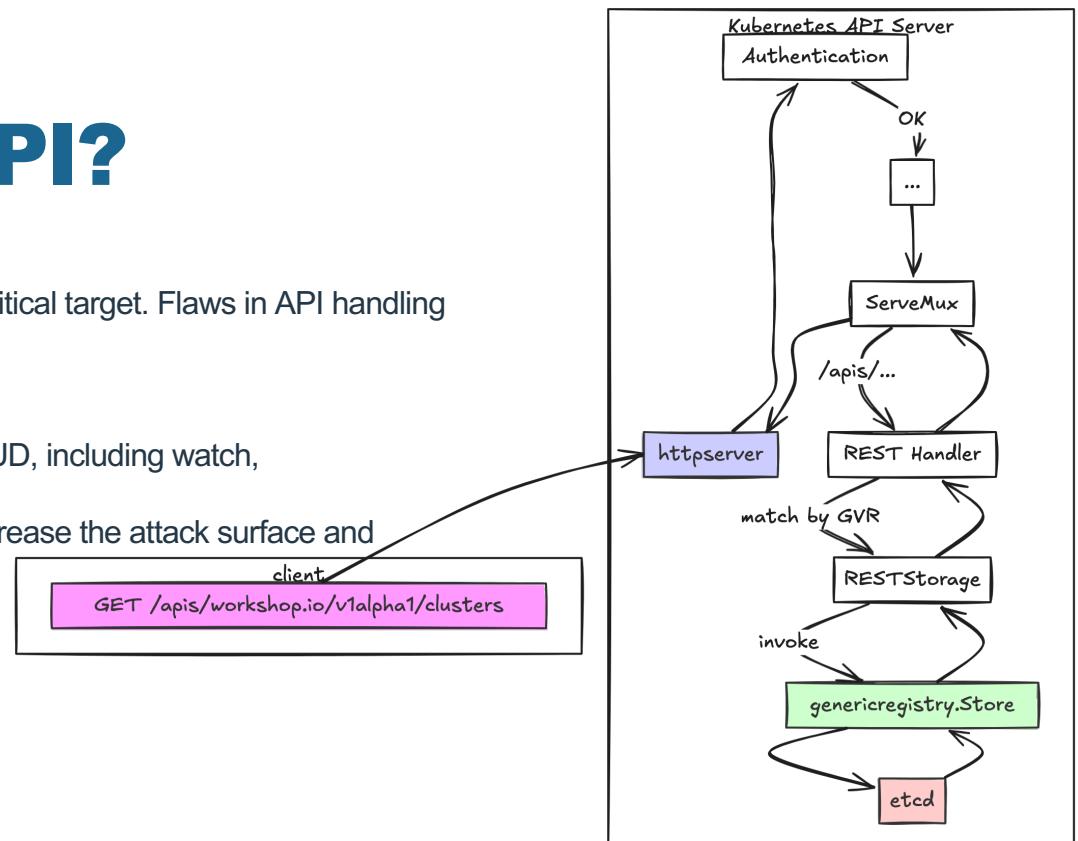
- Built-in sensitive field detection covers all endpoint versions (requires default permissions).
- Match sensitive field patterns by parsing OpenAPI schema

```
4 v var sensitivePatterns = []string{
5     `(?i)password`, // match "password", not case-sensitive
6     `(?i)passwd`,
7     `(?i)pwd`,
8     `(?i)secret`,
9     `(?i)token`,
10    `(?i)api[_-]?key`,
11    `(?i)auth[_-]?token`,
12    `(?i)kubeconfig`,
```



# Why fuzz the k8s API?

- All API resource objects are stored in etcd, making etcd a critical target. Flaws in API handling could expose this data.
- The Kubernetes API supports operations beyond basic CRUD, including watch, deletecollection, and others. These extended operations increase the attack surface and potential for vulnerabilities.



# Understanding Kubernetes API Actions

API Verb	HTTP Method	REST Handler	Description
CREATE	POST	Create(ctx, Obj, ...)	Create new resource
GET	GET	Get(ctx, name, options)	Get a single resource with specified name
LIST	GET	List(ctx, options)	Get resource list
UPDATE	PUT	Update(ctx, name, ...)	Update existing resource
PATCH	PATCH	Update(ctx, name, ...)	Patch a resource
DELETE	DELETE	Delete(ctx, name, ...)	Delete resource
<b>DELETECOLLECTION</b>	DELETE	DeleteCollection(ctx, ..., dryRun)	Delete all resources that match options
<b>WATCH</b>	GET ?watch=true	Watch(ctx, options)	Monitor a resource change in real-time

# Method Promotion: Filter Bypass Risk

- Developers often require finer-grained access control than RBAC provides.
- Developers implement custom filters at the REST layer to enforce these fine-grained policies.
- Developers may **overlook that embedded (anonymous) structs in Go can promote methods**, potentially bypassing these custom filters.

```

type ClusterRest struct {
    *genericregistry.Store
}

func (s *ClusterRest) ShortNames() []string {
    return []string{s.Delete}
}

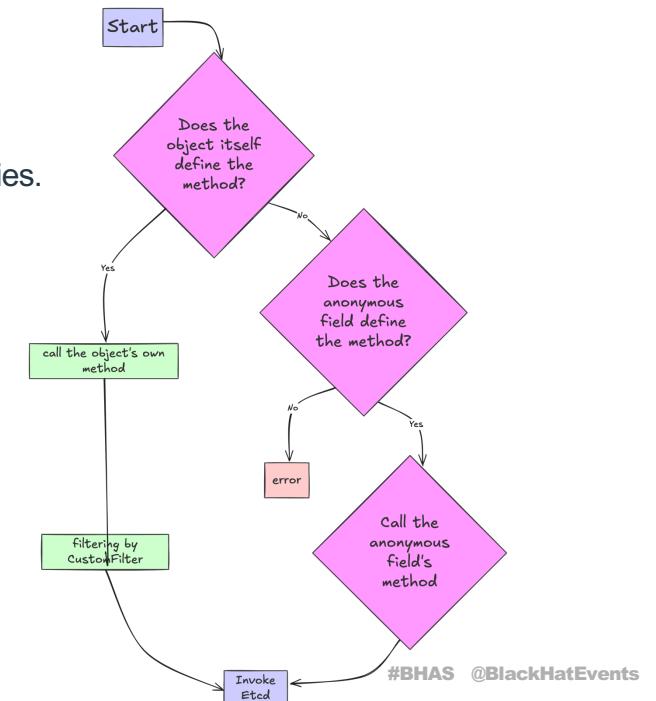
var _ rest.GenericRegistry = (*ClusterRest)(nil)
var _ rest.Store = (*ClusterRest)(nil)

func (s *ClusterRest) Delete(ctx context.Context, name string, deleteValidation rest.ValidateObjectFunc, options *metav1.DeleteOptions) (runtime.Object, bool, error) {
    func (e *genericregistry.Store) DeleteCollection(ctx context.Context, deleteValidation rest.ValidateObjectFunc, options *metav1.DeleteOptions, listOptions *metav1.ListOptions) (runtime.Object, error) {
        return e.DeleteCollection(ctx, deleteValidation, options)
    }
    var _ rest.GenericRegistry = s
    var _ rest.Store = s
    return s.Delete(ctx, name, deleteValidation, options)
}

func (s *ClusterRest) Get(ctx context.Context, name string, options metav1.GetOptions) (runtime.Object, error) {
    return s.GenericRegistry.Get(ctx, name, options)
}

func (s *ClusterRest) GetCreateStrategy() rest.RESTCreateStrategy {
    return s.GenericRegistry.GetCreateStrategy()
}

```



# Simple, But Effective

- The **deletecollection** and **watch** verbs are often overlooked in development.
- fuzzed the API by systematically varying HTTP verbs and API endpoint combinations.
- Use **dryRun** with deletecollection to safely explore its behavior without risking data loss or disruption.
- Use **timeoutSeconds** with watch to control connection duration and prevent resource exhaustion during testing.



# Workshop In Action & Demo

#BHAS @BlackHatEvents



# Workshop: simple multi-tenant cluster

- We'll use a custom resource: workshop.io/v1alpha1/cluster.
- Each Cluster resource includes tenant ownership and kubeconfig information.
- The example include two tenants, tenant-1 and tenant-2, each with their own Cluster

resource, tenant isolation is enforced

```
// Cluster is the Schema for the clusters API
type Cluster struct {
    metav1.TypeMeta `json:",inline"`
    metav1.ObjectMeta `json:"metadata,omitempty" protobuf:"bytes,1,opt,name=metadata"`

    Spec  ClusterSpec `json:"spec,omitempty" protobuf:"bytes,2,opt,name=spec"`
    Status ClusterStatus `json:"status,omitempty"`
}

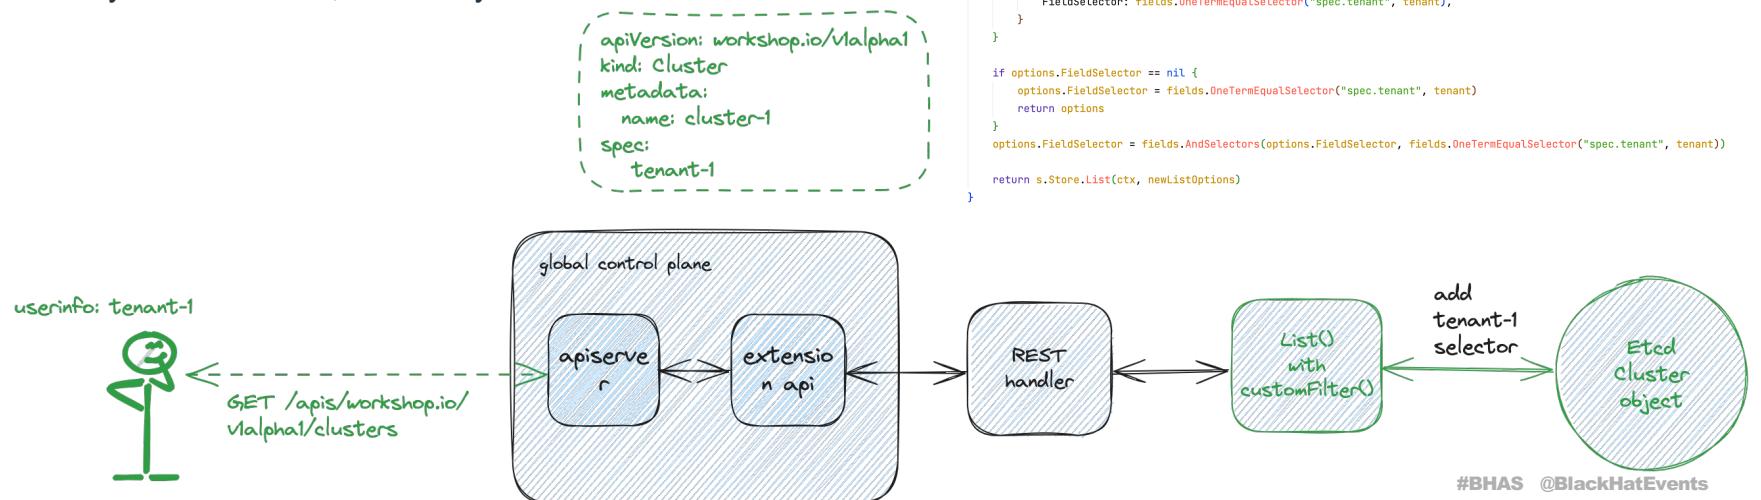
// ClusterSpec defines the desired state of Cluster
type ClusterSpec struct {
    Tenant      string `json:"tenant,omitempty" protobuf:"bytes,1,opt,name=tenant"`
    Name        string `json:"name,omitempty" protobuf:"bytes,2,opt,name=name"`
    APIServer   string `json:"apiserver,omitempty" protobuf:"bytes,3,opt,name=apiserver"`
    Kubeconfig []byte `json:"kubeconfig,omitempty" protobuf:"bytes,4,opt,name=kubeconfig"`
}
```

```
apiVersion: apiregistration.k8s.io/v1
kind: APIService
metadata:
  name: v1alpha1.workshop.io
spec:
  version: v1alpha1
  versionPriority: 1000
  group: workshop.io
  groupPriorityMinimum: 10000
  insecureSkipTLSVerify: true
  service:
    name: workshop-apiservice
    namespace: kubeapi-inspector-workshop
---
apiVersion: v1
kind: Service
metadata:
  name: workshop-apiservice
  namespace: kubeapi-inspector-workshop
  labels:
    app: workshop-apiservice
spec:
  ports:
    - name: apiservice
      port: 443
      protocol: TCP
      targetPort: 443
  selector:
    app: workshop-apiserver
```

#BHAS @BlackHatEvents

# Workshop: simple multi-tenant cluster

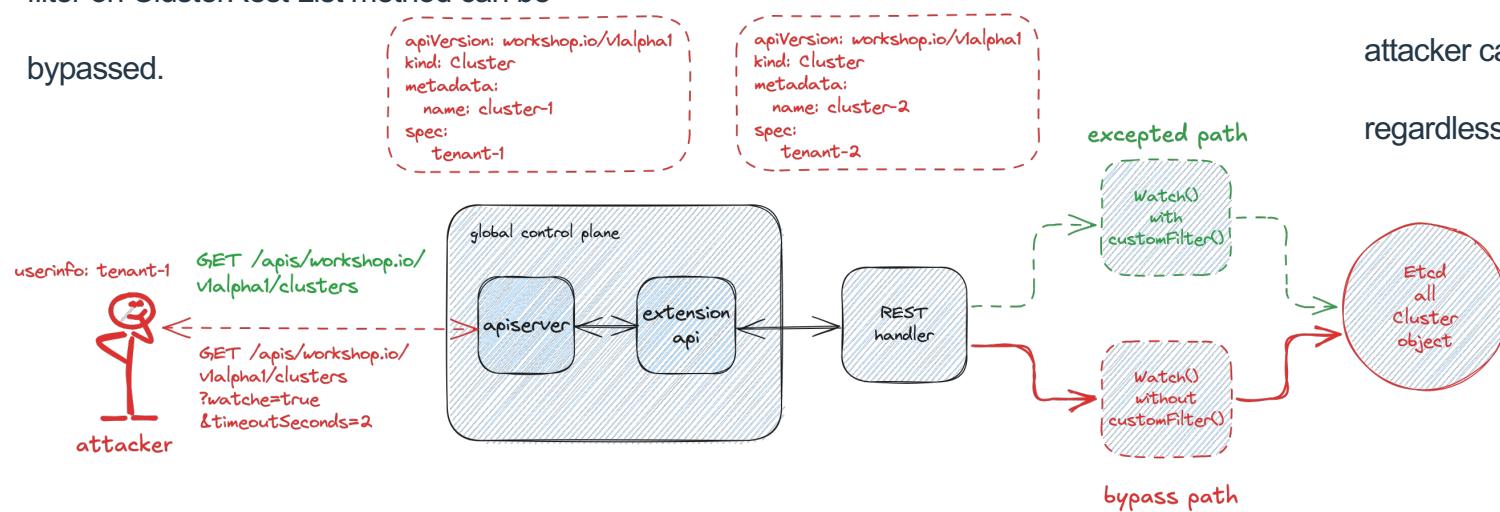
- In the intended setup, a request from tenant-1 for Cluster resources will only return cluster-1, which they own.



#BHAS @BlackHatEvents

# Workshop: Bypass the Resource Filter

- Due to method promotion, the custom filter on ClusterRest List method can be bypassed.



- By crafting a request that targets the embedded Store's List method, an attacker can retrieve all Cluster objects, regardless of tenant.



# Demo

#BHAS @BlackHatEvents

A screenshot of a macOS desktop environment showing a terminal window. The window title is "lazydog@ftw-2:~/workshop". The terminal contains the following command-line session:

```
~ workshop # setup workshop
# minikube start --kubernetes-version='v1.23.17'
~ workshop
```

- Enhance detection of unauthorized access to the Extension API Server
- Expand API fuzzing to include subresource coverage.
- Add more vulnerabilities about kubernetes API in workshop

## Next & Future



<https://github.com/yeahx/KubeAPI-Inspector>



# Thanks

#BHAS @BlackHatEvents